

Machine Learning Assignment 4

Shivangi Aneja

20-November-2017

04_homework_linear_regression

November 19, 2017

1 Programming assignment 4: Linear regression

```
In [180]: import numpy as np
```

```
from sklearn.datasets import load_boston
from sklearn.model_selection import train_test_split
```

1.1 Your task

In this notebook code skeleton for performing linear regression is given. Your task is to complete the functions where required. You are only allowed to use built-in Python functions, as well as any numpy functions. No other libraries / imports are allowed.

1.2 Load and preprocess the data

In this assignment we will work with the Boston Housing Dataset. The data consists of 506 samples. Each sample represents a district in the city of Boston and has 13 features, such as crime rate or taxation level. The regression target is the median house price in the given district (in \$1000's).

More details can be found here: <http://lib.stat.cmu.edu/datasets/boston>

```
In [181]: X , y = load_boston(return_X_y=True)
```

```
# Add a vector of ones to the data matrix to absorb the bias term
# (Recall slide #7 from the lecture)
X = np.hstack([np.ones([X.shape[0], 1]), X])
# From now on, D refers to the number of features in the AUGMENTED dataset (i.e. inclu

# Split into train and test
test_size = 0.2
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=test_size)
```

1.3 Task 1: Fit standard linear regression

```
In [182]: def fit_least_squares(X, y):
    """Fit ordinary least squares model to the data.

    Parameters
```

```

-----
X : array, shape [N, D]
    (Augmented) feature matrix.
y : array, shape [N]
    Regression targets.

Returns
-----
w : array, shape [D]
    Optimal regression coefficients (w[0] is the bias term).

"""
# TODO

X_trans = np.transpose(X)
X_trans_X_inv = np.linalg.inv(np.dot(X_trans,X))
X_trans_y = np.dot(X_trans,y)
weight_array = np.dot(X_trans_X_inv,X_trans_y)

return weight_array

```

1.4 Task 2: Fit ridge regression

```

In [196]: def fit_ridge(X, y, reg_strength):
    """Fit ridge regression model to the data.

    Parameters
    -----
    X : array, shape [N, D]
        (Augmented) feature matrix.
    y : array, shape [N]
        Regression targets.
    reg_strength : float
        L2 regularization strength (denoted by lambda in the lecture)

    Returns
    -----
    w : array, shape [D]
        Optimal regression coefficients (w[0] is the bias term).

    """
```

```

# TODO
X_trans = np.transpose(X)
X_trans_X = np.dot(X_trans,X)
iden = np.identity(X_trans_X.shape[0])
iden_rg = np.multiply(iden,reg_strength)
inv = np.linalg.inv(np.add(X_trans_X,iden_rg))
X_trans_y = np.dot(X_trans,y)

```

```
weight_array = np.dot(inv,X_trans_y)

return weight_array
```

1.5 Task 3: Generate predictions for new data

```
In [172]: def predict_linear_model(X, w):
           """Generate predictions for the given samples.

           Parameters
           -----
           X : array, shape [N, D]
               (Augmented) feature matrix.
           w : array, shape [D]
               Regression coefficients.

           Returns
           -----
           y_pred : array, shape [N]
               Predicted regression targets for the input data.

           """
           # TODO
           X_weighted = X * w
           Y_predicted = X_weighted.sum(axis=1)

           return Y_predicted
```

1.6 Task 4: Mean squared error

```
In [173]: def mean_squared_error(y_true, y_pred):
           """Compute mean squared error between true and predicted regression targets.

           Reference: `https://en.wikipedia.org/wiki/Mean_squared_error`

           Parameters
           -----
           y_true : array
               True regression targets.
           y_pred : array
               Predicted regression targets.

           Returns
           -----
           mse : float
               Mean squared error.

           """
```

```

# TODO
size = y_true.shape[0]
y_diff_sqr = np.square(np.subtract(y_true,y_pred))
y_sum_diff_sqr = np.sum(y_diff_sqr,axis=0)
return np.divide(y_sum_diff_sqr,size)

```

1.7 Compare the two models

The reference implementation produces

* MSE for Least squares ≈ 23.98

* MSE for Ridge regression ≈ 21.05

Your results might be slightly (i.e. $\pm 1\%$) different from the reference solution due to numerical reasons.

```

In [197]: # Load the data
np.random.seed(1234)
X , y = load_boston(return_X_y=True)
X = np.hstack([np.ones([X.shape[0], 1]), X])
test_size = 0.2
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=test_size)

# Ordinary least squares regression
w_ls = fit_least_squares(X_train, y_train)
y_pred_ls = predict_linear_model(X_test, w_ls)
mse_ls = mean_squared_error(y_test, y_pred_ls)
print('MSE for Least squares = {}'.format(mse_ls))

# Ridge regression
reg_strength = 1
w_ridge = fit_ridge(X_train, y_train, reg_strength)
y_pred_ridge = predict_linear_model(X_test, w_ridge)
mse_ridge = mean_squared_error(y_test, y_pred_ridge)
print('MSE for Ridge regression = {}'.format(mse_ridge))

```

MSE for Least squares = 23.9843076118

MSE for Ridge regression = 21.0514870338

Problem 2

Datapoint x_i, y_i weighted by a scalar $t_i > 0$

$$E_{weighted}(w) = \frac{1}{2} \sum_{i=1}^N t_i [w^T \phi(x_i) - y_i]^2 = \frac{1}{2} \sum_{i=1}^N [\sqrt{t_i} (w^T \phi(x_i) - y_i)]^2$$

Converting to vector form , we have

$$\Rightarrow \frac{1}{2} (\phi w - y)^T t (\phi w - y) = \frac{1}{2} (w^T \phi^T - y^T) t (\phi w - y) = \frac{1}{2} (w^T \phi^T t - y^T t) (\phi w - y)$$

$$\Rightarrow \frac{1}{2} [w^T \phi^T t \phi w - y^T t \phi w - w^T \phi^T t y + y^T t y]$$

$$\nabla_w = \phi^T t \phi w - \phi^T t y$$

To find the minimum value, put $\nabla_w = 0$, we have,

$$w^* = (\phi^T t \phi)^{-1} \phi^T t y$$

if $t = I$, then we have, $w_{ML} = (\phi^T \phi)^{-1} \phi^T y$

(1) When we compare w^* to w_{ML} as discusses in the likelihood function of the lecture , then t acts as a precision / inverse Variance for the data point (x_i, y_i)

(2) If $t > 0$ and takes positive Integral values , then more priority to data point with high scaling factor is given for duplicate / replicated data point.

Problem 3

$$E_{ridge}(w) = \frac{1}{2} \sum_{i=1}^N [w^T \phi(x_i) - y_i]^2 + \frac{\lambda}{2} ||w||_2^2$$

$$\phi \in R^{N \times M} \Rightarrow \phi'_{(N+M) \times M} = \begin{bmatrix} \phi \\ \sqrt{\lambda} I_{M \times M} \end{bmatrix}$$

Similarly,

$$y' = \begin{bmatrix} y \\ 0_{M \times 1} \end{bmatrix}$$

We know,

$$w^* = (\phi'^T \phi')^{-1} \phi'^T y \dots\dots(1)$$

$$\phi'^T \phi' = \phi^T \phi + \lambda I \dots\dots(2)$$

$$\phi'^T y' = \phi^T y \dots\dots\dots(3)$$

Using (1),(2) and (3) we have,

$$E_{ridge} = (\phi^T \phi + \lambda I)^{-1} \phi^T y$$

Hence Proved

Problem 4

We have,

$$p(y|\phi, w, \beta) = \prod_{i=1}^N \mathcal{N}(y_i|w^T \phi(x_i), \beta^{-1})$$

$$p(w, \beta) = \mathcal{N}(w|m_0, \beta^{-1} S_0) \text{Gamma}(\beta|a_0, b_0)$$

Posterior \propto Prior \times likelihood

$$p(w, \beta|y) \propto p(y|\phi, w, \beta) p(w|\beta)$$

Taking log on both sides, we have ,

$$\ln p(w, \beta|y) = \ln[p(y|\phi, w, \beta)] + \ln[p(w|\beta)]$$

$$\Rightarrow \ln\left[\prod_{i=1}^N \mathcal{N}(y_i|w^T \phi(x_i), \beta^{-1})\right] + \ln[\mathcal{N}(w|m_0, \beta^{-1} S_0) \text{Gamma}(\beta|a_0, b_0)]$$

$$\Rightarrow \sum_{i=1}^N \ln\left[\frac{1}{\sqrt{2\pi\beta^{-1}}} e^{-\frac{(y_i - w^T \phi(x_i))^2}{2\beta^{-1}}}\right] + \ln\left[\frac{1}{\sqrt{2\pi\beta^{-1} S_0}} e^{-\frac{(w - m_0)^2}{2\beta^{-1} S_0}}\right] + \log\left[\frac{b^{a_0} \beta^{a_0-1} e^{-b_0 \beta}}{\Gamma(a_0)}\right]$$

$$\Rightarrow \sum_{i=1}^N \left(\frac{-1}{2} \ln(2\pi\beta^{-1}) - \frac{(w^T \phi(x_i) - y_i)^2}{2\beta^{-1}}\right) - \frac{1}{2} \ln(2\pi\beta^{-1} S_0) - \frac{(w - m_0)^2}{2\beta^{-1} S_0} + (a_0 - 1) \ln(\beta) - b_0 \beta$$

$$\Rightarrow \frac{N}{2} \ln(\beta) - \frac{\beta}{2} \sum_{i=1}^N (w^T \phi(x_i) - y_i)^2 + (a_0 - 1) \ln(\beta) - b_0 \beta - \frac{\beta}{2} (w - m_0)^T S_0^{-1} (w - m_0) - \frac{1}{2} \ln|S_0| + \frac{m}{2} \ln(\beta)$$

Using Product Rule, we have ,

$$p(w, \beta|y) = p(w|\beta, y) p(\beta|y)$$

$$\ln(p(w|\beta, y)) = \frac{-\beta}{2} w^T (\phi^T \phi + S_0^{-1}) w + w^T [\beta S_0^{-1} m_0 + \beta \phi^T y] + \text{constt.}$$

Thus $p(w|\beta, y)$ is a Gaussian distribution with following mean and covariance

$$m_N = S_N[S_0^{-1}m_0 + \phi^T y]$$

$$S_N^{-1} = \beta(S_0^{-1} + \phi^T \phi)$$

To find $p(\beta|y)$ we need all the terms involvig β and discard terms independent of β

$$\ln(p(\beta|y)) = \frac{-\beta}{2}m_0^T S_0^{-1}m_0 + \frac{\beta}{2}m_N^T S_N^{-1}m_N + \frac{N}{2}\ln(\beta) - b_0\beta + (a_0 - 1)\ln(\beta) - \frac{\beta}{2} \sum_{i=1}^N y_i^2 + \text{constt.}$$

Thus $p(\beta|y)$ is a Gamma distribution with following parameters

$$a_N = a_0 + \frac{N}{2}$$

$$b_N = b_0 + \frac{1}{2}(m_0^T S_0^{-1}m_0 - m_N^T S_N^{-1}m_N + \sum_{i=1}^N y_i^2)$$